

# Deterministic Policy Gradient Algorithms

David Silver, Guy Lever, Nicolas Heess, Thomas Degris, Dann Wierstra, Martin Riedmiller  
Google DeepMind

ICML 2014

Presented by Sanghyeon Lee

# Motivation

## Problem

Previous Policy Gradient methods updating based on action and state

→ *Using Deterministic Policy instead of Stochastic Policy*

1. Propose a Deterministic Policy Algorithm
2. This paper show that Deterministic Policy is the special case of Stochastic Policy
3. DPG is more efficient than SPG
  - DPG has less computation
  - DPG has better performance than SPG especially high dim action space case

# Method

## Notation

$J(\pi) = E[r_1^\gamma | s, a; \pi]$ : Average State value of all state

$\rho^\pi(s) = \lim_{t \rightarrow \infty} P(s_t = s | s_0, \pi_\theta)$  : Stationary distribution of Markov chain for  $\pi_\theta$ , ( $\pi P = \pi$ )

$V^\pi(s) = E[r_1^\gamma | S_1 = s; \pi]$ : Value function, Expected total discounted reward

$Q^\pi(s, a) = E[r_1^\gamma | S_1 = s, A_1 = a; \pi]$       Action Value Function

$\mu_\theta(s)$ : Deterministic Policy

$r_t^\gamma = \sum_{k=t}^{\infty} \gamma^{k-t} r(s_k, a_k)$  : Total Discount Return

$\beta(a|s) \neq \pi_\theta(a|s)$ : Another behavior policy (in off-policy setting)

# Method

## Background

### 1.1 Definition of Value function

- In continuing environments, we can't use discrete value for each state
- Instead of Value function, we can use the average value

$$J_{s_i}(\theta) = V_{\pi_\theta}(s_i) \text{ \& } \pi_\theta(s, a) = P[a|s; \theta]$$

$$J(\pi) = E(r_1^\gamma | \pi) = E_s[V_{\pi_\theta}(s)] = \sum_s \rho^\pi(s) V_{\pi_\theta}(s) = \sum_s \rho^\pi(s) \sum_a \pi_\theta(a|s) Q^\pi(s, a)$$

➔ Average reward per time-step

$$J(\pi_\theta) = \int_S \rho^\pi(s) \int_A \pi_\theta(s, a) r(s, a) da ds = E_{s \sim \rho^\pi, a \sim \pi_\theta}[r(s, a)] \text{ (Integral form)}$$

### 1.2 Stochastic Policy Gradient Theorem

$$\nabla_\theta J(\pi_\theta) = \int_S \rho^\pi(s) \int_A \nabla_\theta \pi_\theta(a|s) Q^\pi(s, a) da ds = E_{s \sim \rho^\pi, a \sim \pi_\theta}[\nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a)]$$

-State Value function has summation form s.t state and action

# Method

## Background

### 1.3 Stochastic Actor-Critic Algorithm

- 1) Critic: Update action-value function
- 2) Actor: Improve policy by Gradient-Descent method

### Action-Value Actor-critic

Simple actor-critic algorithm based on action-value critic

**Input:**  $\pi_\theta, Q_w$ , step size  $\alpha, \beta > 0$

**Initialize**  $s, \theta, w$  at random.

For  $t = 1, \dots, T$

Sample  $R_t \sim r(s, a)$  and the next state  $s' \sim p(s'|s, a)$

Sample next action  $a' \sim \pi_\theta(a'|s')$

$w \leftarrow w + \beta (R_t + \gamma Q_w(s', a') - Q_w(s, a)) \phi(s, a)$  Update critic

$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla_\theta \log \pi_\theta(a|s)$  Update actor

$a \leftarrow a', s \leftarrow s'$

# Method

## Background

### 1.4 Off-Policy Actor-Critic

Off-Policy : Behavior policy  $\neq$  Improvement policy

$$J_{\beta}(\pi_{\theta}) = \int_S \rho^{\beta}(s) V^{\pi}(s) ds = \int_S \int_A \rho^{\beta}(s) \pi_{\theta}(s, a) da ds$$

$$\nabla_{\theta} J_{\beta}(\pi_{\theta}) \approx \int_S \int_A \rho^{\beta}(s) \nabla \pi_{\theta}(a|s) Q^{\pi}(s, a) da ds = E_{s \sim \rho^{\beta}, a \sim \beta} \left[ \frac{\pi_{\theta}(a|s)}{\beta_{\theta}(a|s)} \nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi}(s, a) \right]$$

Pf)  $\nabla_{\theta} J_{\beta}(\pi_{\theta}) = \int_S \int_A \rho^{\beta}(s) [\nabla \pi_{\theta}(a|s) Q^{\pi}(s, a) + \pi_{\theta}(a|s) \nabla Q^{\pi}(s, a)] da ds$   
 $\approx \int_S \int_A \rho^{\beta}(s) \nabla \pi_{\theta}(a|s) Q^{\pi}(s, a) da ds$  \*Degris, 2012b

- Importance Sampling

$$E_{x \sim p} = E_{x \sim q} \left[ \frac{p}{q} f(x) \right]$$

---

**Algorithm 1** The Off-PAC algorithm

---

Initialize the vectors  $\mathbf{e}_v$ ,  $\mathbf{e}_u$ , and  $\mathbf{w}$  to zero

Initialize the vectors  $\mathbf{v}$  and  $\mathbf{u}$  arbitrarily

Initialize the state  $s$

For each step:

Choose an action,  $a$ , according to  $b(\cdot|s)$

Observe resultant reward,  $r$ , and next state,  $s'$

$$\delta \leftarrow r + \gamma(s') \mathbf{v}^T \mathbf{x}_{s'} - \mathbf{v}^T \mathbf{x}_s$$

$$\rho \leftarrow \pi_{\mathbf{u}}(a|s) / b(a|s)$$

Update the critic (GTD( $\lambda$ ) algorithm):

$$\mathbf{e}_v \leftarrow \rho (\mathbf{x}_s + \gamma(s) \lambda \mathbf{e}_v)$$

$$\mathbf{v} \leftarrow \mathbf{v} + \alpha_v [\delta \mathbf{e}_v - \gamma(s') (1 - \lambda) (\mathbf{w}^T \mathbf{e}_v) \mathbf{x}_s]$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha_w [\delta \mathbf{e}_v - (\mathbf{w}^T \mathbf{x}_s) \mathbf{x}_s]$$

Update the actor:

$$\mathbf{e}_u \leftarrow \rho \left[ \frac{\nabla_{\mathbf{u}} \pi_{\mathbf{u}}(a|s)}{\pi_{\mathbf{u}}(a|s)} + \gamma(s) \lambda \mathbf{e}_u \right]$$

$$\mathbf{u} \leftarrow \mathbf{u} + \alpha_u \delta \mathbf{e}_u$$

$$s \leftarrow s'$$

---

# Method

## Gradient of Deterministic Policies

### 2.1 Action-Value Gradients

- Model free RL use greedy policy (Greedy policy:  $\mu^{k+1}(s) = \operatorname{argmax}_a Q^{\mu^k}(s, a)$ )
- In the continuous action space, greedy policy improvement needs a global maximization at every step (It needs large cost)
- Instead of greedy policy, we can improve policy by maximize action-value function  $Q^{\mu^k}(s, a)$

$$\theta^{k+1} = \theta^k + \alpha \mathbb{E}_{s \sim \rho^{\mu^k}} [\nabla_{\theta} Q^{\mu^k}(s, \mu_{\theta}(s))] = \theta^k + \alpha \mathbb{E}_{s \sim \rho^{\mu^k}} [\nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu^k}(s, a)|_{a=\mu_{\theta}(s)}] \text{ (Chain rule)}$$

### 2.2 Deterministic Policy Gradient Theorem

$$J(\mu_{\theta}) = \int_S \rho^{\mu}(s) r(s, \mu_{\theta}(s)) ds = \mathbb{E}_{s \sim \rho^{\mu}} [r(s, \mu_{\theta}(s))]$$

$$\nabla_{\theta} J(\mu_{\theta}) = \int_S \rho^{\mu}(s) \nabla_{\theta}(s) \nabla_a Q^{\mu}(s, a) \Big|_{a=\mu_{\theta}(s)} ds = \mathbb{E}_{s \sim \rho^{\mu}} [\nabla_{\theta} \mu_{\theta}(s) \nabla_a Q^{\mu}(s, a) \Big|_{a=\mu_{\theta}(s)}]$$

- Deterministic policy only needs state distribution  $\rho^{\mu}$  (SPG need  $\rho^{\mu}$  and action distribution)

### 2.3 Limit of the Stochastic Policy Gradient = Deterministic Policy Gradient

$$\lim_{\sigma \downarrow 0} \nabla_{\theta} J(\pi_{\mu_{\theta}, \sigma}) = \nabla_{\theta} J(\mu_{\theta})$$

→ Deterministic policy gradients are familiar of policy gradients

**Regularity conditions A.1:**  $p(s'|s, a)$ ,  $\nabla_a p(s'|s, a)$ ,  $\mu_{\theta}(s)$ ,  $\nabla_{\theta} \mu_{\theta}(s)$ ,  $r(s, a)$ ,  $\nabla_a r(s, a)$ ,  $p_1(s)$  are continuous in all parameters and variables  $s$ ,  $a$ ,  $s'$  and  $x$ .

→ Lipschitz continuous & boundary condition

**Regularity conditions A.2:** there exists a  $b$  and  $L$  such that  $\sup_s p_1(s) < b$ ,  $\sup_{a, s, s'} p(s'|s, a) < b$ ,  $\sup_{a, s} r(s, a) < b$ ,  $\sup_{a, s, s'} \|\nabla_a p(s'|s, a)\| < L$ , and  $\sup_{a, s} \|\nabla_a r(s, a)\| < L$ .

# Method

## Apply DPG – Deterministic Actor-Critic Algorithms

### 3.1 On-Policy Deterministic Actor-Critic

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, a_{t+1}) - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t) \quad (\nabla_w Q^w = w^T \nabla \phi(s, a))$$

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu_\theta(s_t) \nabla_a Q^w(s_t, a_t)|_{a=\mu_\theta(s)} \rightarrow \text{Only Actor change}$$

### 3.2 Off-Policy Deterministic Actor-Critic (OPDAC)

-We can avoid importance sampling in the actor and by using Q-learning, we can also avoid importance sampling in the critic

$$J_\beta(\mu_\theta) = \int_s \rho^\beta V^\mu(s) ds = \int_s \rho^\beta(s) Q^\mu(s, \mu_\theta(s)) ds$$

$$\nabla_\theta J_\beta(\mu_\theta) \approx \int_s \rho^\beta(s) \nabla_\theta \mu_\theta(a|s) Q^\mu(s, a) ds = \mathbb{E}_{s \sim \rho^\beta} [\nabla_\theta \mu_\theta(s) \nabla_a Q^\mu(s, a)|_{a=\mu_\theta(s)}]$$

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, a_{t+1}) - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t)$$

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu_\theta(s_t) \nabla_a Q^w(s_t, a_t)|_{a=\mu_\theta(s)}$$



# Method

## Apply DPG – Deterministic Actor–Critic Algorithms

### 3.3 Compatible Function Approximation

– An approximator  $Q^w(s, a)$  is not necessarily compatible with true gradient ;  $Q^\mu(s, a)$

Theorem 3.

1)  $\nabla_a Q^w(s, a)|_{a=\mu_\theta(s)} = \nabla_\theta \mu_\theta(s)^T w$  & 2)  $MSE(\theta, w) = E[\epsilon(s; \theta, w)^T \epsilon(s; \theta, w)]$  where

$$\epsilon(s; \theta, w) = \nabla_a Q^w(s, a) \Big|_{a=\mu_\theta(s)} - \nabla_a Q^\mu(s, a) \Big|_{a=\mu_\theta(s)}$$

Then, a function approximator  $Q^w(s, a)$  is compatible with a deterministic policy  $\mu_\theta(s)$ ,

$$\nabla_\theta J_\beta(\theta) = E[\nabla_\theta \mu_\theta(s) \nabla_a Q^w(s, a) \Big|_{a=\mu_\theta(s)}]$$

### Reduce Variance

– We can express approximator action–value function with baseline function independent of the action  $a$

$$Q^w(s, a) = (a - \mu_\theta(s))^T \nabla_\theta \mu_\theta(s)^T w + V^v(s), V^v \text{ is independent with action}$$

(ex  $V^v(s) = v^T \phi(s)$  for parameters  $v$ )

- If  $a \approx \mu_\theta$ :  $Q^w(s, a) \approx V^v(s)$
- $A^w(s, a) = \phi(s, a)^T w$ ,  $\phi(s, a) = \nabla_\theta \mu_\theta(s)(a - \mu_\theta(s))$
- Advantage function  $A^w(s, a) = Q_w(s, a) - V(s)$ , ex)  $V(s) \approx \frac{1}{N} \sum Q^w(s_{i,t}, a_{i,t})$

# Method

## Apply DPG – Deterministic Actor–Critic Algorithms

Compatible off-policy deterministic actor critic with Q-learning (COPDAC-Q)

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, \mu_\theta(s_{t+1})) - Q^w(s_t, a_t)$$

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu_\theta(s_t) (\nabla_\theta \mu_\theta(s_t)^\top w_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \phi(s_t, a_t)$$

$$v_{t+1} = v_t + \alpha_v \delta_t \phi(s_t)$$

### Theorem3

$\nabla_\theta J_\beta(\theta) = E[\nabla_\theta \mu_\theta(s) \nabla_a Q^w(s, a)|_{a=\mu_\theta(s)}]$ ,  
V is independent with action

# Method

\* Equivalent forms of policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) G_t] \quad \text{REINFORCE}$$

$$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q_w(s, a)] \quad \text{Q Actor-critic}$$

$$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) A_w(s, a)] \quad \text{Advantage Actor-critic}$$

$$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) \delta_t] \quad \text{TD Actor-critic}$$

$$\delta_v = r + V_{\xi}(s') - V_{\xi}(s)$$

# Experiments

## 1. Continuous Bandit

- Multidimensional bandit problem ( $a \in R^m$ ),  
*Regret( difference between optimal policy rewards and choosen policy*  
$$\text{Regret}(-r(a)) = (a - a^*)^T C(a - a^*),$$
$$C: \text{psd with eigenvalue} \in \{0.1, 1\}, a^* = [4, \dots, 4]^T \in R^m)$$
- SAC-B:  $\pi_{\theta,y}(\cdot) \sim N(\theta, \exp(y))$ , COPDAC-B:  $\mu_\theta = \theta$  (mean of Gaussian)
- Critic is computed from each successive batch of 2m steps





Figure 1. Comparison of stochastic actor-critic (SAC-B) and deterministic actor-critic (COPDAC-B) on the continuous bandit task.


# Experiments

## 2. Continuous Reinforcement Learning (mountain car, pendulum and 2D puddle world)


- SAC-Q:  $\pi_{\theta,y}(\cdot) \sim N(\theta^T \phi(s), \exp(y^T \phi(s)))$ ,  $V(s) = v^T \phi(s)$
- COPDAC-Q:  $\mu_{\theta}(s) = \theta^T \phi(s)$ ,  $\beta(\cdot | s) \sim N(\theta^T \phi(s), \sigma_{\beta}^2)$ ,  $V(s) = v^T \phi(s)$
- OffPAC-TD:  $\beta(\cdot | s) \sim N(\theta^T \phi(s), \sigma_{\beta}^2)$ ,  $\pi_{\theta,y}(\cdot) \sim N(\theta^T \phi(s), \exp(y^T \phi(s)))$ ,  $V(s) = v^T \phi(s)$



(a) Mountain Car



(b) Pendulum




(c) 2D Puddle World

Figure 2. Comparison of stochastic on-policy actor-critic (SAC), stochastic off-policy actor-critic (OffPAC), and deterministic off-policy actor-critic (COPDAC) on continuous-action reinforcement learning. Each point is the average test performance of the mean policy.

# Experiments

## 3. Octopus Arm



*Figure 3.* Ten runs of COPDAC on a 6-segment octopus arm with 20 action dimensions and 50 state dimensions; each point represents the return per episode (above) and the number of time-steps for the arm to reach the target (below).

# Discussion

1. DPG is more efficient than SPG
  - DPG has less computation
  - DPG has better performance than SPG especially high dim action space case
2. But, It is not suitable in the game which has stochastic optimal policy